

# Exam2-Part2-Boston housing

Take home (65 pts) Read the question carefully and write your answers briefly supporting your conclusions with plots and statistical quantities.

You need to submit the html file and the Rmd file. Before submitting verify your html file that it includes all necessary plots and also check all necessary values are printed in the html.

Note that the questions below are open ended. There is no fixed “correct” answer. Try to use various ideas and techniques we have seen during the class.

The data set Boston contains data on Boston housing prices. The data consist of the 506 houses in Boston area.

The response variable is

- $Y = \text{medv}$  = median value of owner-occupied homes in \$1000s.

The predictor variables are:

- $\text{crim}$  = per capita crime rate by town,
- $\text{zn}$  = proportion of residential land zoned for lots over 25,000 sq.ft.,
- $\text{indus}$  = proportion of non-retail business acres per town,
- $\text{chas}$  = Charles River dummy variable (= 1 if tract bounds river; 0 otherwise)
- $\text{nox}$  = nitrogen oxides concentration (parts per 10 million)
- $\text{rm}$  = average number of rooms per dwelling,
- $\text{age}$  = proportion of owner-occupied units built prior to 1940,
- $\text{dis}$  = weighted mean of distances to five Boston employment centers
- $\text{rad}$  = index of accessibility to radial highways
- $\text{tax}$  = full-value property-tax rate per \$10,000
- $\text{pratio}$  = pupil-teacher ratio by town
- $\text{lstat}$  = lower status of the population (percent). Your goal is to develop a model that predicts median value of the house ( $\text{medv}$ ). You start with the multiple linear regression model using all of the 12 regressors (this is your base model). Answer the below questions. In all parts write your model clearly. In addition to writing your justification clearly, print the critical values and display the plots you use.

- Base model:** Fit a multiple linear regression model using all 12 regressors.

Answer goes here (model and summary):

```
library(readxl)
Boston <- read_excel("Downloads/Boston_housing.xlsx")

Boston.BM<- lm(medv~crim+zn+indus+as.factor(chas)+nox+rm+age+dis+as.factor(rad)+tax+ptra
tio+lstat, data = Boston)
summary(Boston.BM)
```

```
##
## Call:
## lm(formula = medv ~ crim + zn + indus + as.factor(chas) + nox +
##      rm + age + dis + as.factor(rad) + tax + ptratio + lstat,
##      data = Boston)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -14.6357  -2.7013  -0.5723   1.8160  25.9979
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   40.398739   5.296438   7.628 1.27e-13 ***
## crim          -0.121816   0.032858  -3.707 0.000234 ***
## zn             0.055525   0.014314   3.879 0.000119 ***
## indus          0.016795   0.064363   0.261 0.794250
## as.factor(chas)1  2.677692   0.872194   3.070 0.002260 **
## nox          -18.455862   3.933930  -4.691 3.53e-06 ***
## rm             3.511231   0.423837   8.284 1.16e-15 ***
## age            0.003511   0.013353   0.263 0.792741
## dis          -1.568899   0.204235  -7.682 8.72e-14 ***
## as.factor(rad)2   1.527760   1.494794   1.022 0.307264
## as.factor(rad)3   4.698681   1.350945   3.478 0.000550 ***
## as.factor(rad)4   2.606331   1.201262   2.170 0.030516 *
## as.factor(rad)5   2.864862   1.221675   2.345 0.019427 *
## as.factor(rad)6   1.283888   1.480915   0.867 0.386394
## as.factor(rad)7   4.917263   1.589585   3.093 0.002093 **
## as.factor(rad)8   4.820869   1.509140   3.194 0.001492 **
## as.factor(rad)24  7.123585   1.807059   3.942 9.26e-05 ***
## tax           -0.009111   0.003939  -2.313 0.021146 *
## ptratio       -0.960781   0.146134  -6.575 1.26e-10 ***
## lstat        -0.557596   0.050584 -11.023 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 4.749 on 486 degrees of freedom
## Multiple R-squared:  0.7434, Adjusted R-squared:  0.7334
## F-statistic: 74.12 on 19 and 486 DF, p-value: < 2.2e-16
```

```
attach(Boston)
```

- b. Interaction Terms:** Give a model that uses base model and includes interaction terms Crim x age, rm x tax, rm x ptratio, tax x ptratio, nox x crim, nox x age and 3 additional interaction terms of your choice. Check if any of these interaction terms contribute to the model. Do backwards selection to create a simpler model with interactions.

Answer goes here (model and summary):

```
Boston.IT<- lm(medv~crim+zn+indus+as.factor(chas)+nox+rm+age+dis+as.factor(rad)+tax+ptratio+lstat+crim*age+rm*tax+rm*ptratio+tax*ptratio+nox*crim+nox*age+indus*tax+crim*tax+ptratio*crim, data = Boston)
summary(Boston.IT)
```

```
##
## Call:
## lm(formula = medv ~ crim + zn + indus + as.factor(chas) + nox +
##      rm + age + dis + as.factor(rad) + tax + ptratio + lstat +
##      crim * age + rm * tax + rm * ptratio + tax * ptratio + nox *
##      crim + nox * age + indus * tax + crim * tax + ptratio * crim,
##      data = Boston)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -10.9424  -2.2374  -0.3788   1.3927  26.8693
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   -9.635e+01  2.268e+01  -4.247 2.60e-05 ***
## crim           5.849e+00  2.807e+00   2.084 0.037679 *
## zn            2.534e-02  1.396e-02   1.816 0.070027 .
## indus         2.188e-01  1.332e-01   1.643 0.101125
## as.factor(chas)1 3.571e+00  7.586e-01   4.707 3.31e-06 ***
## nox           6.504e-02  1.296e+01   0.005 0.995998
## rm            2.274e+01  2.406e+00   9.452 < 2e-16 ***
## age           4.881e-02  6.902e-02   0.707 0.479760
## dis          -8.554e-01  1.904e-01  -4.494 8.79e-06 ***
## as.factor(rad)2  1.561e+00  1.310e+00   1.192 0.234000
## as.factor(rad)3  4.707e+00  1.187e+00   3.964 8.49e-05 ***
## as.factor(rad)4  2.016e+00  1.073e+00   1.879 0.060817 .
## as.factor(rad)5  2.499e+00  1.072e+00   2.330 0.020218 *
## as.factor(rad)6  1.912e+00  1.308e+00   1.462 0.144408
## as.factor(rad)7  3.874e+00  1.383e+00   2.801 0.005300 **
## as.factor(rad)8  3.272e+00  1.329e+00   2.461 0.014196 *
## as.factor(rad)24 6.252e+00  1.778e+00   3.516 0.000479 ***
## tax           6.487e-02  3.815e-02   1.700 0.089715 .
## ptratio       3.543e+00  1.305e+00   2.714 0.006883 **
## lstat        -5.098e-01  4.603e-02 -11.075 < 2e-16 ***
## crim:age       6.644e-03  3.281e-03   2.025 0.043413 *
## rm:tax        -1.429e-02  2.234e-03  -6.395 3.82e-10 ***
## rm:ptratio    -6.943e-01  1.557e-01  -4.459 1.03e-05 ***
## tax:ptratio    1.061e-03  1.986e-03   0.534 0.593387
## crim:nox      -1.303e+00  6.620e-01  -1.968 0.049611 *
## nox:age       -1.398e-01  1.420e-01  -0.984 0.325547
## indus:tax     -3.664e-04  3.594e-04  -1.020 0.308475
## crim:tax      -1.922e-03  2.929e-03  -0.656 0.511955
## crim:ptratio  -2.219e-01  1.916e-01  -1.158 0.247438
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 4.098 on 477 degrees of freedom
## Multiple R-squared:  0.8125, Adjusted R-squared:  0.8015
## F-statistic: 73.82 on 28 and 477 DF,  p-value: < 2.2e-16
```

```
Boston.IT2= update(Boston.IT,~.- nox)
summary(Boston.IT2)
```

```
##
## Call:
## lm(formula = medv ~ crim + zn + indus + as.factor(chas) + rm +
##      age + dis + as.factor(rad) + tax + ptratio + lstat + crim:age +
##      rm:tax + rm:ptratio + tax:ptratio + crim:nox + nox:age +
##      indus:tax + crim:tax + crim:ptratio, data = Boston)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -10.9395  -2.2382  -0.3797   1.3924  26.8683
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   -9.631e+01  2.183e+01  -4.413 1.26e-05 ***
## crim           5.850e+00  2.800e+00   2.090 0.037177 *
## zn            2.531e-02  1.295e-02   1.955 0.051217 .
## indus         2.187e-01  1.323e-01   1.653 0.099088 .
## as.factor(chas)1 3.571e+00  7.573e-01   4.715 3.18e-06 ***
## rm            2.274e+01  2.401e+00   9.471 < 2e-16 ***
## age           4.850e-02  3.003e-02   1.615 0.106989
## dis          -8.557e-01  1.830e-01  -4.676 3.82e-06 ***
## as.factor(rad)2  1.561e+00  1.308e+00   1.193 0.233491
## as.factor(rad)3  4.707e+00  1.183e+00   3.980 7.98e-05 ***
## as.factor(rad)4  2.016e+00  1.071e+00   1.882 0.060430 .
## as.factor(rad)5  2.500e+00  1.066e+00   2.344 0.019496 *
## as.factor(rad)6  1.913e+00  1.295e+00   1.477 0.140467
## as.factor(rad)7  3.874e+00  1.381e+00   2.804 0.005250 **
## as.factor(rad)8  3.272e+00  1.325e+00   2.470 0.013862 *
## as.factor(rad)24 6.253e+00  1.767e+00   3.539 0.000441 ***
## tax           6.485e-02  3.795e-02   1.709 0.088139 .
## ptratio       3.543e+00  1.302e+00   2.722 0.006720 **
## lstat        -5.098e-01  4.588e-02 -11.113 < 2e-16 ***
## crim:age       6.637e-03  2.964e-03   2.239 0.025585 *
## rm:tax        -1.428e-02  2.217e-03  -6.442 2.89e-10 ***
## rm:ptratio    -6.943e-01  1.550e-01  -4.481 9.31e-06 ***
## tax:ptratio    1.062e-03  1.983e-03   0.535 0.592668
## crim:nox      -1.303e+00  6.609e-01  -1.971 0.049253 *
## age:nox       -1.391e-01  5.127e-02  -2.714 0.006898 **
## indus:tax     -3.663e-04  3.584e-04  -1.022 0.307246
## crim:tax      -1.924e-03  2.917e-03  -0.660 0.509869
## crim:ptratio  -2.218e-01  1.913e-01  -1.160 0.246778
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 4.093 on 478 degrees of freedom
## Multiple R-squared:  0.8125, Adjusted R-squared:  0.8019
## F-statistic: 76.72 on 27 and 478 DF,  p-value: < 2.2e-16
```

```
Boston.IT3= update(Boston.IT2,~.-tax*ptratio)
summary(Boston.IT3)
```

```
##
## Call:
## lm(formula = medv ~ crim + zn + indus + as.factor(chas) + rm +
##      age + dis + as.factor(rad) + lstat + crim:age + rm:tax +
##      rm:ptratio + crim:nox + age:nox + indus:tax + crim:tax +
##      crim:ptratio, data = Boston)
##
## Residuals:
##      Min        1Q    Median        3Q        Max
## -10.590   -2.571   -0.478    1.790   25.679
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    4.3100153   3.4575740     1.247 0.213172
## crim           0.0276931   2.5699636     0.011 0.991407
## zn             0.0474615   0.0137271     3.458 0.000594 ***
## indus        -0.2743951   0.1117556    -2.455 0.014429 *
## as.factor(chas)1  2.8210865   0.8151229     3.461 0.000586 ***
## rm            8.8659180   0.6052396    14.649 < 2e-16 ***
## age           0.0607075   0.0323404     1.877 0.061104 .
## dis          -1.1493713   0.1943650    -5.913 6.36e-09 ***
## as.factor(rad)2   1.8944467   1.4138169     1.340 0.180894
## as.factor(rad)3   3.8763580   1.2770792     3.035 0.002533 **
## as.factor(rad)4   2.1460206   1.1580557     1.853 0.064477 .
## as.factor(rad)5   2.8963317   1.1513826     2.516 0.012210 *
## as.factor(rad)6   2.6115996   1.4009698     1.864 0.062911 .
## as.factor(rad)7   4.5657020   1.4913667     3.061 0.002326 **
## as.factor(rad)8   4.4139516   1.4223641     3.103 0.002027 **
## as.factor(rad)24  9.1352099   1.7170950     5.320 1.59e-07 ***
## lstat          -0.5319757   0.0487698   -10.908 < 2e-16 ***
## crim:age         0.0084197   0.0031944     2.636 0.008665 **
## rm:tax           -0.0055165   0.0008270    -6.670 7.04e-11 ***
## rm:ptratio       -0.1592776   0.0249827    -6.376 4.28e-10 ***
## crim:nox         -1.3107301   0.7154998    -1.832 0.067583 .
## age:nox          -0.1580175   0.0552165    -2.862 0.004396 **
## indus:tax         0.0011841   0.0002703     4.380 1.46e-05 ***
## crim:tax         -0.0054103   0.0030043    -1.801 0.072354 .
## crim:ptratio      0.1741724   0.1829360     0.952 0.341527
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 4.435 on 481 degrees of freedom
## Multiple R-squared:  0.7785, Adjusted R-squared:  0.7674
## F-statistic: 70.44 on 24 and 481 DF, p-value: < 2.2e-16
```

```
Boston.ITF= update(Boston.IT3,~.-crim)
summary(Boston.ITF)
```

```
##
## Call:
## lm(formula = medv ~ zn + indus + as.factor(chas) + rm + age +
##      dis + as.factor(rad) + lstat + crim:age + rm:tax + rm:ptratio +
##      crim:nox + age:nox + indus:tax + crim:tax + crim:ptratio,
##      data = Boston)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -10.5895  -2.5698  -0.4786   1.7872  25.6796
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    4.3122447   3.4477964    1.251 0.211642
## zn              0.0474543   0.0136970    3.465 0.000578 ***
## indus          -0.2741699   0.1096703   -2.500 0.012752 *
## as.factor(chas)1  2.8213666   0.8138626    3.467 0.000574 ***
## rm              8.8677469   0.5803503   15.280 < 2e-16 ***
## age             0.0606034   0.0308342    1.965 0.049934 *
## dis            -1.1492385   0.1937725   -5.931 5.75e-09 ***
## as.factor(rad)2   1.8937441   1.4108470    1.342 0.180139
## as.factor(rad)3   3.8776503   1.2701162    3.053 0.002391 **
## as.factor(rad)4   2.1447319   1.1506684    1.864 0.062943 .
## as.factor(rad)5   2.8975848   1.1443052    2.532 0.011652 *
## as.factor(rad)6   2.6110856   1.3987045    1.867 0.062538 .
## as.factor(rad)7   4.5666931   1.4869831    3.071 0.002253 **
## as.factor(rad)8   4.4148538   1.4184242    3.113 0.001965 **
## as.factor(rad)24  9.1322636   1.6934265    5.393 1.09e-07 ***
## lstat           -0.5320396   0.0483578  -11.002 < 2e-16 ***
## age:crim         0.0084210   0.0031889    2.641 0.008542 **
## rm:tax           -0.0055158   0.0008238   -6.696 5.99e-11 ***
## rm:ptratio       -0.1594187   0.0212591   -7.499 3.12e-13 ***
## crim:nox         -1.3102825   0.7135518   -1.836 0.066932 .
## age:nox          -0.1578292   0.0523218   -3.017 0.002692 **
## indus:tax         0.0011840   0.0002699    4.387 1.41e-05 ***
## crim:tax         -0.0054212   0.0028269   -1.918 0.055737 .
## crim:ptratio      0.1758814   0.0910785    1.931 0.054057 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 4.431 on 482 degrees of freedom
## Multiple R-squared:  0.7785, Adjusted R-squared:  0.7679
## F-statistic: 73.65 on 23 and 482 DF,  p-value: < 2.2e-16
```


c. **Transformation of variables:** Try various transformations on the base model, then propose a transformation (on prediction variable medv or on regressors) that you think it might be helpful to linearize the model (or to improve it). Then fit a model using this transformation. Explain which variables were transformed and why.

Answer goes here (model, summary and explanation):


```
library(car)
```

```
## Loading required package: carData
```

```
plot(crim~medv)
```




```
plot(lstat~medv)
```




```
boxCox(lm(crim~medv), family= "yjPower")
```





```
boxCox(lm(lstat~medv), family= "yjPower")
```



```
crimT<- yjPower(crim, -.9)
lstatT<- yjPower(lstat, -.1)
Boston.TM<- lm(medv~crimT+zn+indus+as.factor(chas)+nox+rm+age+dis+as.factor(rad)+tax+ptratio+lstatT, data = Boston)
summary(Boston.TM)
```

```
##
## Call:
## lm(formula = medv ~ crimT + zn + indus + as.factor(chas) + nox +
##      rm + age + dis + as.factor(rad) + tax + ptratio + lstatT,
##      data = Boston)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -14.0180  -2.5715  -0.2805   1.9719  25.5426
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    64.316215    5.319075   12.092 < 2e-16 ***
## crimT         -0.111819    1.708098   -0.065 0.947832
## zn             0.025549    0.013179    1.939 0.053136 .
## indus          0.027352    0.059198    0.462 0.644258
## as.factor(chas)1  2.418201    0.802432    3.014 0.002717 **
## nox          -15.645542    3.935633   -3.975 8.09e-05 ***
## rm             2.345261    0.404626    5.796 1.22e-08 ***
## age            0.029704    0.012687    2.341 0.019618 *
## dis          -1.184617    0.187578   -6.315 6.09e-10 ***
## as.factor(rad)2  1.445469    1.372572    1.053 0.292814
## as.factor(rad)3  4.214801    1.242199    3.393 0.000748 ***
## as.factor(rad)4  2.579657    1.126363    2.290 0.022434 *
## as.factor(rad)5  2.569778    1.129175    2.276 0.023292 *
## as.factor(rad)6  2.408636    1.362104    1.768 0.077635 .
## as.factor(rad)7  4.601021    1.473805    3.122 0.001904 **
## as.factor(rad)8  3.878574    1.416535    2.738 0.006407 **
## as.factor(rad)24  5.749277    1.984332    2.897 0.003933 **
## tax           -0.009116    0.003617   -2.520 0.012043 *
## ptratio       -0.858215    0.137905   -6.223 1.05e-09 ***
## lstatT       -13.767180    0.833010  -16.527 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 4.36 on 486 degrees of freedom
## Multiple R-squared:  0.7837, Adjusted R-squared:  0.7752
## F-statistic: 92.67 on 19 and 486 DF,  p-value: < 2.2e-16
```

```
#I will try another transformation this time only transforming the response variable
Boston.TM2<- lm(log(medv)~crim+zn+indus+as.factor(chas)+nox+rm+age+dis+as.factor(rad)+tax+ptratio+lstat, data = Boston)
summary(Boston.TM2)
```

```
##
## Call:
## lm(formula = log(medv) ~ crim + zn + indus + as.factor(chas) +
##      nox + rm + age + dis + as.factor(rad) + tax + ptratio + lstat,
##      data = Boston)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.68178 -0.10160 -0.01198  0.09992  0.81278
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   4.2347902   0.2138514   19.802 < 2e-16 ***
## crim         -0.0108009   0.0013267   -8.141 3.32e-15 ***
## zn            0.0015374   0.0005779    2.660 0.00807 **
## indus         0.0024672   0.0025988    0.949 0.34290
## as.factor(chas)1 0.1070532   0.0352161    3.040 0.00249 **
## nox          -0.8118549   0.1588382   -5.111 4.61e-07 ***
## rm            0.0800587   0.0171130    4.678 3.76e-06 ***
## age           0.0003328   0.0005392    0.617 0.53734
## dis          -0.0517614   0.0082463   -6.277 7.66e-10 ***
## as.factor(rad)2  0.0850869   0.0603545    1.410 0.15924
## as.factor(rad)3  0.1774313   0.0545464    3.253 0.00122 **
## as.factor(rad)4  0.1015640   0.0485027    2.094 0.03678 *
## as.factor(rad)5  0.1321144   0.0493269    2.678 0.00765 **
## as.factor(rad)6  0.1004904   0.0597941    1.681 0.09348 .
## as.factor(rad)7  0.2092091   0.0641818    3.260 0.00119 **
## as.factor(rad)8  0.1931184   0.0609337    3.169 0.00162 **
## as.factor(rad)24 0.3355303   0.0729627    4.599 5.43e-06 ***
## tax          -0.0005212   0.0001590   -3.277 0.00112 **
## ptratio       -0.0364930   0.0059004   -6.185 1.32e-09 ***
## lstat        -0.0304173   0.0020424  -14.893 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.1917 on 486 degrees of freedom
## Multiple R-squared:  0.7883, Adjusted R-squared:  0.78
## F-statistic: 95.22 on 19 and 486 DF, p-value: < 2.2e-16
```

### ***#What if we only transform one predictor variable***

```
Boston.TM3<- lm(medv~crim+zn+indus+as.factor(chas)+nox+rm+age+dis+as.factor(rad)+tax+ptratio+lstatT, data = Boston)
summary(Boston.TM3)
```

```
##
## Call:
## lm(formula = medv ~ crim + zn + indus + as.factor(chas) + nox +
##      rm + age + dis + as.factor(rad) + tax + ptratio + lstatT,
##      data = Boston)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -14.4979  -2.5389  -0.2708   1.8781  25.0625
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   65.601308   5.203532  12.607 < 2e-16 ***
## crim         -0.140215   0.029157  -4.809 2.03e-06 ***
## zn            0.030528   0.012885   2.369 0.018210 *
## indus         0.017518   0.057689   0.304 0.761519
## as.factor(chas)1  2.197908   0.783715   2.804 0.005242 **
## nox          -16.553602   3.531636  -4.687 3.60e-06 ***
## rm            2.249945   0.395187   5.693 2.16e-08 ***
## age           0.030719   0.012212   2.515 0.012209 *
## dis          -1.289352   0.184539  -6.987 9.30e-12 ***
## as.factor(rad)2   1.327724   1.340175   0.991 0.322321
## as.factor(rad)3   4.194652   1.210190   3.466 0.000575 ***
## as.factor(rad)4   2.545643   1.076744   2.364 0.018461 *
## as.factor(rad)5   2.558322   1.096226   2.334 0.020015 *
## as.factor(rad)6   2.231204   1.331308   1.676 0.094391 .
## as.factor(rad)7   4.715777   1.424782   3.310 0.001003 **
## as.factor(rad)8   3.994674   1.353551   2.951 0.003318 **
## as.factor(rad)24  7.325010   1.620564   4.520 7.77e-06 ***
## tax           -0.009265   0.003534  -2.622 0.009020 **
## ptratio       -0.882820   0.131163  -6.731 4.77e-11 ***
## lstatT        -13.402069   0.817409 -16.396 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 4.26 on 486 degrees of freedom
## Multiple R-squared:  0.7935, Adjusted R-squared:  0.7854
## F-statistic:  98.3 on 19 and 486 DF,  p-value: < 2.2e-16
```

**#Notice we have achived a higher R2 when only lstat is transformed. This transformation might be helpful to linearize the model**

```
bsell=step(Boston.TM3)
```

```
## Start:  AIC=1486.3
## medv ~ crim + zn + indus + as.factor(chas) + nox + rm + age +
##       dis + as.factor(rad) + tax + ptratio + lstatT
##
##              Df Sum of Sq    RSS    AIC
## - indus       1      1.7  8822.2 1484.4
## <none>                8820.5 1486.3
## - zn          1    101.9  8922.4 1490.1
## - age         1    114.8  8935.3 1490.8
## - tax         1    124.8  8945.3 1491.4
## - as.factor(chas) 1    142.7  8963.2 1492.4
## - nox         1    398.7  9219.2 1506.7
## - as.factor(rad)  8    669.7  9490.2 1507.3
## - crim        1    419.7  9240.2 1507.8
## - rm          1    588.3  9408.8 1517.0
## - ptratio     1    822.2  9642.7 1529.4
## - dis         1    886.0  9706.5 1532.7
## - lstatT      1   4878.9 13699.4 1707.1
##
## Step:  AIC=1484.39
## medv ~ crim + zn + as.factor(chas) + nox + rm + age + dis + as.factor(rad) +
##       tax + ptratio + lstatT
##
##              Df Sum of Sq    RSS    AIC
## <none>                8822.2 1484.4
## - zn          1    100.2  8922.4 1488.1
## - age         1    114.8  8937.0 1488.9
## - tax         1    137.8  8959.9 1490.2
## - as.factor(chas) 1    147.5  8969.7 1490.8
## - nox         1    417.3  9239.5 1505.8
## - crim        1    422.0  9244.2 1506.0
## - as.factor(rad)  8    692.0  9514.2 1506.6
## - rm          1    588.0  9410.1 1515.0
## - ptratio     1    820.5  9642.7 1527.4
## - dis         1    936.9  9759.1 1533.5
## - lstatT      1   4882.1 13704.3 1705.3
```

```
library(MASS)
```

```
##
## Attaching package: 'MASS'
```

```
## The following object is masked _by_ '.GlobalEnv':
##
##      Boston
```

```
step1 <- stepAIC(Boston.TM3, direction="both")
```

```
## Start:  AIC=1486.3
## medv ~ crim + zn + indus + as.factor(chas) + nox + rm + age +
##      dis + as.factor(rad) + tax + ptratio + lstatT
##
##              Df Sum of Sq    RSS    AIC
## - indus      1      1.7  8822.2 1484.4
## <none>                        8820.5 1486.3
## - zn         1    101.9  8922.4 1490.1
## - age        1    114.8  8935.3 1490.8
## - tax        1    124.8  8945.3 1491.4
## - as.factor(chas) 1    142.7  8963.2 1492.4
## - nox        1    398.7  9219.2 1506.7
## - as.factor(rad)  8    669.7  9490.2 1507.3
## - crim       1    419.7  9240.2 1507.8
## - rm         1    588.3  9408.8 1517.0
## - ptratio    1    822.2  9642.7 1529.4
## - dis        1    886.0  9706.5 1532.7
## - lstatT     1   4878.9 13699.4 1707.1
##
## Step:  AIC=1484.39
## medv ~ crim + zn + as.factor(chas) + nox + rm + age + dis + as.factor(rad) +
##      tax + ptratio + lstatT
##
##              Df Sum of Sq    RSS    AIC
## <none>                        8822.2 1484.4
## + indus      1      1.7  8820.5 1486.3
## - zn         1    100.2  8922.4 1488.1
## - age        1    114.8  8937.0 1488.9
## - tax        1    137.8  8959.9 1490.2
## - as.factor(chas) 1    147.5  8969.7 1490.8
## - nox        1    417.3  9239.5 1505.8
## - crim       1    422.0  9244.2 1506.0
## - as.factor(rad)  8    692.0  9514.2 1506.6
## - rm         1    588.0  9410.1 1515.0
## - ptratio    1    820.5  9642.7 1527.4
## - dis        1    936.9  9759.1 1533.5
## - lstatT     1   4882.1 13704.3 1705.3
```

```
step1$anova # display results
```

```
## Stepwise Model Path
## Analysis of Deviance Table
##
## Initial Model:
## medv ~ crim + zn + indus + as.factor(chas) + nox + rm + age +
##      dis + as.factor(rad) + tax + ptratio + lstatT
##
## Final Model:
## medv ~ crim + zn + as.factor(chas) + nox + rm + age + dis + as.factor(rad) +
##      tax + ptratio + lstatT
##
##
```

	Step	Df	Deviance	Resid. Df	Resid. Dev	AIC
## 1				486	8820.493	1486.298
## 2 - indus	1	1.67349		487	8822.167	1484.394

***#Notice that the backwards stepping and stepwise stepping both found the same model.***

```
BestTMM<- lm(medv ~ crim + zn + as.factor(chas) + nox + rm + age + dis + as.factor(rad)
+
            tax + ptratio + lstatT, data = Boston)
```

I choose to transform Crim and Lstat because when I looked at the relation of them to medv they both were non linear.

- d. **Polynomial terms:** Eliminate 6 of the regressors from the base model, that (you think) are the least significant ones. (You can do a subjective choice, considering the nature of the data, as long as you support it. For example you can make a few joint significance test to support your choice). Now using the remaining 6 regressors propose a polynomial model that includes quadratic terms and interaction terms. Then fit this model. Answer goes here (model, summary and explanation):

```
bsel2<- step(Boston.BM)
```



```
## Start:  AIC=1596.16
## medv ~ crim + zn + indus + as.factor(chas) + nox + rm + age +
##      dis + as.factor(rad) + tax + ptratio + lstat
##
##              Df Sum of Sq  RSS    AIC
## - indus      1      1.54 10961 1594.2
## - age        1      1.56 10961 1594.2
## <none>                      10959 1596.2
## - tax        1    120.63 11080 1599.7
## - as.factor(chas) 1    212.54 11172 1603.9
## - crim       1    309.94 11269 1608.3
## - zn         1    339.33 11299 1609.6
## - nox        1    496.32 11456 1616.6
## - as.factor(rad)  8    820.77 11780 1616.7
## - ptratio    1    974.75 11934 1637.3
## - dis        1   1330.69 12290 1652.1
## - rm         1   1547.64 12507 1661.0
## - lstat      1   2740.02 13699 1707.1
##
## Step:  AIC=1594.23
## medv ~ crim + zn + as.factor(chas) + nox + rm + age + dis + as.factor(rad) +
##      tax + ptratio + lstat
##
##              Df Sum of Sq  RSS    AIC
## - age        1      1.55 10962 1592.3
## <none>                      10961 1594.2
## - tax        1    133.52 11094 1598.4
## - as.factor(chas) 1    218.62 11180 1602.2
## - crim       1    312.24 11273 1606.4
## - zn         1    340.09 11301 1607.7
## - as.factor(rad)  8    843.21 11804 1615.7
## - nox        1    521.81 11483 1615.8
## - ptratio    1    973.36 11934 1635.3
## - dis        1   1400.69 12362 1653.1
## - rm         1   1553.55 12514 1659.3
## - lstat      1   2743.42 13704 1705.3
##
## Step:  AIC=1592.3
## medv ~ crim + zn + as.factor(chas) + nox + rm + dis + as.factor(rad) +
##      tax + ptratio + lstat
##
##              Df Sum of Sq  RSS    AIC
## <none>                      10962 1592.3
## - tax        1    132.35 11095 1596.4
## - as.factor(chas) 1    221.02 11184 1600.4
## - crim       1    312.36 11275 1604.5
## - zn         1    338.89 11301 1605.7
## - as.factor(rad)  8    841.91 11804 1613.7
## - nox        1    543.93 11506 1614.8
## - ptratio    1    975.77 11938 1633.5
## - dis        1   1563.90 12526 1657.8
## - rm         1   1623.80 12586 1660.2
## - lstat      1   3050.72 14013 1714.5
```

```
summary(bsel2)
```

```
##
## Call:
## lm(formula = medv ~ crim + zn + as.factor(chas) + nox + rm +
##     dis + as.factor(rad) + tax + ptratio + lstat, data = Boston)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -14.6779  -2.7212  -0.4832   1.7849  26.1280
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    40.237252     5.267117   7.639 1.16e-13 ***
## crim          -0.122189     0.032768  -3.729 0.000215 ***
## zn              0.054603     0.014058   3.884 0.000117 ***
## as.factor(chas)1  2.712564     0.864787   3.137 0.001812 **
## nox          -17.911598     3.640033  -4.921 1.18e-06 ***
## rm              3.519984     0.414017   8.502 2.29e-16 ***
## dis          -1.594663     0.191121  -8.344 7.44e-16 ***
## as.factor(rad)2   1.597403     1.476432   1.082 0.279816
## as.factor(rad)3   4.674248     1.346722   3.471 0.000565 ***
## as.factor(rad)4   2.616563     1.195405   2.189 0.029081 *
## as.factor(rad)5   2.852057     1.218321   2.341 0.019635 *
## as.factor(rad)6   1.222933     1.467568   0.833 0.405080
## as.factor(rad)7   4.897766     1.585188   3.090 0.002118 **
## as.factor(rad)8   4.805755     1.498994   3.206 0.001434 **
## as.factor(rad)24  6.997177     1.767284   3.959 8.64e-05 ***
## tax           -0.008623     0.003552  -2.427 0.015572 *
## ptratio        -0.954861     0.144881  -6.591 1.14e-10 ***
## lstat         -0.552240     0.047388 -11.654 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 4.74 on 488 degrees of freedom
## Multiple R-squared:  0.7434, Adjusted R-squared:  0.7344
## F-statistic: 83.15 on 17 and 488 DF,  p-value: < 2.2e-16
```

```
#I am going to remove age and indus from our base model since from backwards selection w
e notice that it is not significant
# I will also remove chas, dis, and tax as I do not believe that they would influence the
medv
Boston.P4Update<- update(bsel2,~.-as.factor(chas)-tax-dis)
summary(Boston.P4Update)
```

```
##
## Call:
## lm(formula = medv ~ crim + zn + nox + rm + as.factor(rad) + ptratio +
##      lstat, data = Boston)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -14.2080  -3.1390  -0.7746   1.9087  28.9452
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    21.511053     5.160991   4.168 3.63e-05 ***
## crim          -0.096564     0.035215  -2.742 0.006327 **
## zn            -0.007202     0.013196  -0.546 0.585453
## nox          -3.983067     3.323917  -1.198 0.231376
## rm             4.292130     0.438902   9.779 < 2e-16 ***
## as.factor(rad)2  3.144138     1.580501   1.989 0.047219 *
## as.factor(rad)3  5.290236     1.450650   3.647 0.000294 ***
## as.factor(rad)4  3.166493     1.285905   2.462 0.014141 *
## as.factor(rad)5  2.980062     1.315260   2.266 0.023901 *
## as.factor(rad)6  1.497927     1.560891   0.960 0.337698
## as.factor(rad)7  3.426066     1.707781   2.006 0.045388 *
## as.factor(rad)8  4.971560     1.617701   3.073 0.002235 **
## as.factor(rad)24 5.018556     1.471749   3.410 0.000703 ***
## ptratio       -1.092511     0.155188  -7.040 6.52e-12 ***
## lstat         -0.537326     0.051082 -10.519 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 5.134 on 491 degrees of freedom
## Multiple R-squared:  0.697, Adjusted R-squared:  0.6884
## F-statistic: 80.69 on 14 and 491 DF, p-value: < 2.2e-16
```

**# And from a backwards selection of the Boston.P4update I will remove zn from our model since it is also not significant**

```
Boston.P4Update2<- update(Boston.P4Update,~.-zn)
summary(Boston.P4Update2)
```

```
##
## Call:
## lm(formula = medv ~ crim + nox + rm + as.factor(rad) + ptratio +
##      lstat, data = Boston)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -14.190  -3.077  -0.792   1.895  28.975
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    20.37551     4.71969   4.317 1.91e-05 ***
## crim          -0.09693     0.03518  -2.755 0.006087 **
## nox           -3.22554     3.01812  -1.069 0.285718
## rm             4.27712     0.43773   9.771 < 2e-16 ***
## as.factor(rad)2  3.27544     1.56097   2.098 0.036385 *
## as.factor(rad)3  5.44426     1.42192   3.829 0.000145 ***
## as.factor(rad)4  3.25228     1.27535   2.550 0.011071 *
## as.factor(rad)5  3.13387     1.28380   2.441 0.014995 *
## as.factor(rad)6  1.63192     1.54036   1.059 0.289921
## as.factor(rad)7  3.51028     1.69958   2.065 0.039411 *
## as.factor(rad)8  5.18316     1.56943   3.303 0.001028 **
## as.factor(rad)24 5.04332     1.47000   3.431 0.000652 ***
## ptratio        -1.05905     0.14246  -7.434 4.71e-13 ***
## lstat          -0.53623     0.05101 -10.513 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 5.13 on 492 degrees of freedom
## Multiple R-squared:  0.6968, Adjusted R-squared:  0.6888
## F-statistic: 86.99 on 13 and 492 DF,  p-value: < 2.2e-16
```

***#Now I will start to add polynomial terms. First I would like to look at correlations.***

```
cor(crim,crim^2)
```

```
## [1] 0.8710611
```

```
cor(nox,nox^2)
```

```
## [1] 0.9935007
```

```
cor(lstat,lstat^2)
```

```
## [1] 0.9605726
```

```
cor(rm,rm^2)
```

```
## [1] 0.994528
```

```
cor(ptratio,ptratio^2)
```

```
## [1] 0.9979917
```

```
#the correlation between all of the above variables and their square could be a problem  
#lets try the square of the transformed variables  
tcrim<-(Boston$crim-mean(Boston$crim))/sd(Boston$crim)  
tlstat<-(Boston$lstat-mean(Boston$lstat))/sd(Boston$lstat)
```

```
cor(crim,tcrim^2)
```

```
## [1] 0.8365473
```

```
cor(lstat,tlstat^2)
```

```
## [1] 0.5742952
```

```
Boston.Poly<- update(Boston.P4Update2,~.+I(tlstat^2)+I(tcrim^2)+crim*age+rm*tax+rm*ptratio+nox*crim+nox*age+indus*tax+crim*tax+ptratio*crim, data = Boston)  
summary(Boston.Poly)
```

```
##
## Call:
## lm(formula = medv ~ crim + nox + rm + as.factor(rad) + ptratio +
##      lstat + I(tlstat^2) + I(tcrim^2) + age + tax + indus + crim:age +
##      rm:tax + rm:ptratio + crim:nox + nox:age + tax:indus + crim:tax +
##      crim:ptratio, data = Boston)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -9.7767 -2.4045 -0.3254  1.7169 26.8505
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   -1.079e+02  1.602e+01  -6.738 4.62e-11 ***
## crim           4.437e+00  2.491e+00   1.781 0.075499 .
## nox            2.007e+01  1.066e+01   1.883 0.060293 .
## rm             2.187e+01  2.253e+00   9.703 < 2e-16 ***
## as.factor(rad)2  1.290e+00  1.268e+00   1.018 0.309395
## as.factor(rad)3  4.954e+00  1.139e+00   4.349 1.67e-05 ***
## as.factor(rad)4  2.632e+00  1.041e+00   2.528 0.011806 *
## as.factor(rad)5  2.407e+00  1.037e+00   2.321 0.020730 *
## as.factor(rad)6  2.947e+00  1.270e+00   2.320 0.020755 *
## as.factor(rad)7  3.780e+00  1.341e+00   2.819 0.005011 **
## as.factor(rad)8  3.483e+00  1.278e+00   2.725 0.006656 **
## as.factor(rad)24 9.190e+00  1.692e+00   5.431 8.93e-08 ***
## ptratio        4.248e+00  9.832e-01   4.320 1.90e-05 ***
## lstat         -7.789e-01  5.852e-02 -13.310 < 2e-16 ***
## I(tlstat^2)     1.446e+00  1.817e-01   7.957 1.29e-14 ***
## I(tcrim^2)      3.035e-01  8.347e-02   3.636 0.000307 ***
## age            1.458e-01  5.775e-02   2.525 0.011906 *
## tax             5.397e-02  1.685e-02   3.204 0.001448 **
## indus          2.565e-01  1.164e-01   2.204 0.028024 *
## crim:age        4.782e-03  3.175e-03   1.506 0.132693
## rm:tax          -9.858e-03  2.339e-03  -4.215 2.99e-05 ***
## rm:ptratio      -7.737e-01  1.473e-01  -5.254 2.25e-07 ***
## crim:nox        -2.053e+00  6.431e-01  -3.192 0.001505 **
## nox:age         -2.575e-01  1.211e-01  -2.126 0.034015 *
## tax:indus       -2.877e-04  3.036e-04  -0.947 0.343910
## crim:tax        2.733e-04  2.850e-03   0.096 0.923647
## crim:ptratio    -2.042e-01  1.798e-01  -1.136 0.256573
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 4 on 479 degrees of freedom
## Multiple R-squared:  0.8206, Adjusted R-squared:  0.8109
## F-statistic: 84.28 on 26 and 479 DF,  p-value: < 2.2e-16
```

***#Now I will do a step wise to find the best model.***

```
step3 <- stepAIC(Boston.Poly, direction="both")
```

```

## Start:  AIC=1429.1
## medv ~ crim + nox + rm + as.factor(rad) + ptratio + lstat + I(tlstat^2) +
##      I(tcrim^2) + age + tax + indus + crim:age + rm:tax + rm:ptratio +
##      crim:nox + nox:age + tax:indus + crim:tax + crim:ptratio
##
##              Df Sum of Sq    RSS    AIC
## - crim:tax      1      0.15  7662.9 1427.1
## - tax:indus      1     14.36  7677.1 1428.0
## - crim:ptratio   1     20.64  7683.4 1428.5
## <none>                        7662.8 1429.1
## - crim:age       1     36.29  7699.1 1429.5
## - nox:age        1     72.31  7735.1 1431.8
## - crim:nox       1    163.01  7825.8 1437.8
## - I(tcrim^2)     1    211.48  7874.3 1440.9
## - rm:tax         1    284.22  7947.0 1445.5
## - rm:ptratio     1    441.53  8104.3 1455.5
## - as.factor(rad)  8    813.00  8475.8 1464.1
## - I(tlstat^2)    1   1012.76  8675.5 1489.9
## - lstat          1   2834.24 10497.0 1586.3
##
## Step:  AIC=1427.11
## medv ~ crim + nox + rm + as.factor(rad) + ptratio + lstat + I(tlstat^2) +
##      I(tcrim^2) + age + tax + indus + crim:age + rm:tax + rm:ptratio +
##      crim:nox + nox:age + tax:indus + crim:ptratio
##
##              Df Sum of Sq    RSS    AIC
## - tax:indus      1     14.22  7677.2 1426.0
## <none>                        7662.9 1427.1
## - crim:age       1     37.12  7700.1 1427.6
## - crim:ptratio   1     42.87  7705.8 1427.9
## + crim:tax       1      0.15  7662.8 1429.1
## - nox:age        1     72.16  7735.1 1429.9
## - crim:nox       1    166.90  7829.8 1436.0
## - I(tcrim^2)     1    211.77  7874.7 1438.9
## - rm:tax         1    297.82  7960.8 1444.4
## - rm:ptratio     1    466.64  8129.6 1455.0
## - as.factor(rad)  8    832.13  8495.1 1463.3
## - I(tlstat^2)    1   1015.19  8678.1 1488.1
## - lstat          1   2844.73 10507.7 1584.9
##
## Step:  AIC=1426.05
## medv ~ crim + nox + rm + as.factor(rad) + ptratio + lstat + I(tlstat^2) +
##      I(tcrim^2) + age + tax + indus + crim:age + rm:tax + rm:ptratio +
##      crim:nox + nox:age + crim:ptratio
##
##              Df Sum of Sq    RSS    AIC
## <none>                        7677.2 1426.0
## - crim:age       1     37.26  7714.4 1426.5
## - crim:ptratio   1     45.42  7722.6 1427.0
## + tax:indus      1     14.22  7662.9 1427.1
## + crim:tax       1      0.01  7677.1 1428.0
## - nox:age        1     68.43  7745.6 1428.5
## - indus          1    138.61  7815.8 1433.1

```

```
## - crim:nox      1      167.04  7844.2 1434.9
## - I(tcrim^2)    1      205.74  7882.9 1437.4
## - rm:tax        1      285.99  7963.1 1442.6
## - rm:ptratio    1      489.34  8166.5 1455.3
## - as.factor(rad) 8      821.20  8498.4 1461.5
## - I(tlstat^2)   1     1017.94  8695.1 1487.0
## - lstat         1     2862.08 10539.2 1584.4
```

```
step3$anova # display results
```

```
## Stepwise Model Path
## Analysis of Deviance Table
##
## Initial Model:
## medv ~ crim + nox + rm + as.factor(rad) + ptratio + lstat + I(tlstat^2) +
##      I(tcrim^2) + age + tax + indus + crim:age + rm:tax + rm:ptratio +
##      crim:nox + nox:age + tax:indus + crim:tax + crim:ptratio
##
## Final Model:
## medv ~ crim + nox + rm + as.factor(rad) + ptratio + lstat + I(tlstat^2) +
##      I(tcrim^2) + age + tax + indus + crim:age + rm:tax + rm:ptratio +
##      crim:nox + nox:age + crim:ptratio
##
##
##          Step Df   Deviance Resid. Df Resid. Dev      AIC
## 1              479   7662.785 1429.103
## 2 - crim:tax    1   0.1471001   480   7662.932 1427.112
## 3 - tax:indus   1 14.2218170   481   7677.154 1426.051
```

```
BestPY<- lm(medv ~ crim + nox + rm + as.factor(rad) + ptratio + lstat + I(tlstat^2) +
            I(tcrim^2) + age + tax + indus + crim*age + rm*tax + rm*ptratio +
            crim*nox + nox*age + crim*ptratio, data = Boston)
```

e. **Compare** the performance of the models in part (a) to (d). Look at various diagnostics we have seen. Also check for normality and constant variance violations. Make a comparison and support your comments with plots and statistics.

Answer goes here:

```
print(paste0("Boston Base Model adjusted R squared    ", summary(Boston.BM)$adj.r.squared
, " Boston Base Model residual standard error      ", summary(Boston.BM)$sigma) )
```

```
## [1] "Boston Base Model adjusted R squared    0.733408239842023 Boston Base Model resi
dual standard error    4.7486970348662"
```

```
print(paste0("Boston Interaction Model adjusted R squared    ", summary(Boston.ITF)$adj.
r.squared," Boston Interaction Model residual standard error    ", summary(Boston.ITF)
$sigma) )
```



```
## [1] "Boston Interaction Model adjusted R squared    0.767925400503753  Boston Interact  
ion Model residual standard error    4.43062410962611"
```





```
print(paste0("Boston Polynomial Model adjusted R squared    ", summary(BestPY)$adj.r.squa  
red, "  Boston Polynomial Model residual standard error    ",  summary(BestPY)$sigma))
```

```
## [1] "Boston Polynomial Model adjusted R squared    0.811308222633533  Boston Polynomia  
l Model residual standard error    3.99509940008713"
```

```
print(paste0("Boston Transfromation Model adjusted R squared    ", summary(BestTMM)$adj.  
r.squared, "  Boston Transformation Model residual standard error    ",  summary(BestTM  
M)$sigma) )
```

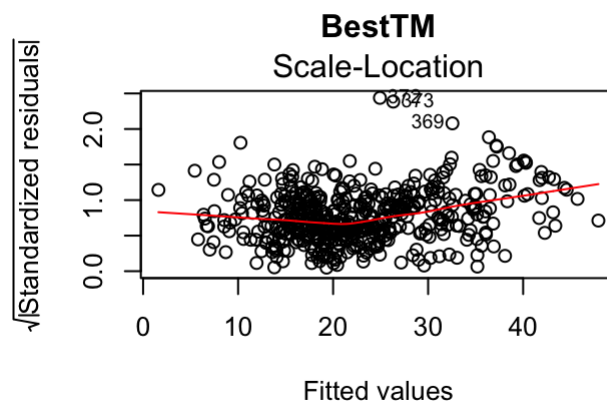
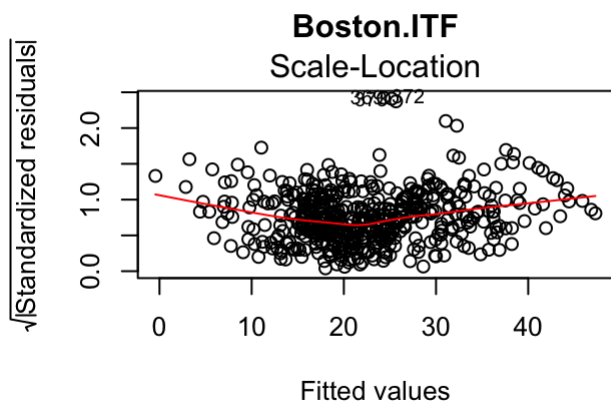

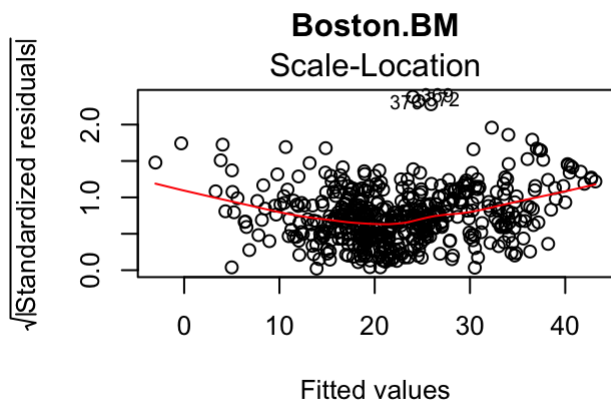
```
## [1] "Boston Transfromation Model adjusted R squared    0.785837164391323  Boston Trans  
formation Model residual standard error    4.25621105911814"
```

```
#testing for normaility  
layout(matrix(c(1,2,3,4),2,2))  
qqnorm(residuals(Boston.ITF), main='Boston.ITF')  
qqline(residuals(Boston.ITF), col='red')  
qqnorm(residuals(BestPY), main='BestPY')  
qqline(residuals(BestPY), col='red')  
qqnorm(residuals(BestTMM), main='BestTMM')  
qqline(residuals(BestTMM), col='red')  
qqnorm(residuals(Boston.BM), main='Boston.BM')  
qqline(residuals(Boston.BM), col='red')
```

**Boston.ITF****BestTMM****BestPY****Boston.BM**

***#Looking at Standardized residuals to determine if there is a constant variance***

```
layout(matrix(c(1,2,3,4),2,2))
plot(Boston.BM,3, main='Boston.BM')
plot(Boston.ITF, 3, main='Boston.ITF')
plot(BestPY, 3, main='BestPY')
plot(BestTMM,3, main='BestTM')
```



From

the above comparisons we can see that the best fitted model is The Boston Polynomial Model (BestPY). It is the best fitted model since it has the highest adjusted  $R^2$  and standard error. It is also the most normal model and seems to have a constant variance.

- f. **Make your own:** Now considering all of the above, propose a new model different than the ones in part a-d (try mixture of the suggestions above). Use best subsets to fit your model. Comment on overall adequacy of your model comparing with the ones above.

Answer goes here (model, summary, explanation and comparison):

```
library(leaps)
model.subset <- regsubsets(log(medv)~crim+zn+indus+as.factor(chas)+nox+rm+age+dis+as.factor(rad)+tax+prratio+lstat+ I(tlstat^2) + crim*age + rm*tax + rm*prratio + crim*nox + nox*age + crim*prratio, data = Boston, nbest = 1, nvmax = 26)
summary(model.subset)
```

```
## Subset selection object
## Call: regsubsets.formula(log(medv) ~ crim + zn + indus + as.factor(chas) +
##      nox + rm + age + dis + as.factor(rad) + tax + ptratio + lstat +
##      I(tlstat^2) + crim * age + rm * tax + rm * ptratio + crim *
##      nox + nox * age + crim * ptratio, data = Boston, nbest = 1,
##      nvmax = 26)
## 26 Variables (and intercept)
##
```

	Forced in	Forced out
## crim	FALSE	FALSE
## zn	FALSE	FALSE
## indus	FALSE	FALSE
## as.factor(chas)1	FALSE	FALSE
## nox	FALSE	FALSE
## rm	FALSE	FALSE
## age	FALSE	FALSE
## dis	FALSE	FALSE
## as.factor(rad)2	FALSE	FALSE
## as.factor(rad)3	FALSE	FALSE
## as.factor(rad)4	FALSE	FALSE
## as.factor(rad)5	FALSE	FALSE
## as.factor(rad)6	FALSE	FALSE
## as.factor(rad)7	FALSE	FALSE
## as.factor(rad)8	FALSE	FALSE
## as.factor(rad)24	FALSE	FALSE
## tax	FALSE	FALSE
## ptratio	FALSE	FALSE
## lstat	FALSE	FALSE
## I(tlstat^2)	FALSE	FALSE
## crim:age	FALSE	FALSE
## rm:tax	FALSE	FALSE
## rm:ptratio	FALSE	FALSE
## crim:nox	FALSE	FALSE
## nox:age	FALSE	FALSE
## crim:ptratio	FALSE	FALSE

```
## 1 subsets of each size up to 26
## Selection Algorithm: exhaustive
##
```

	crim	zn	indus	as.factor(chas)1	nox	rm	age	dis	as.factor(rad)2
## 1 ( 1 )	" "	" "	" "	" "	" "	" "	" "	" "	" "
## 2 ( 1 )	" "	" "	" "	" "	" "	" "	" "	" "	" "
## 3 ( 1 )	" "	" "	" "	" "	" "	" "	" "	" "	" "
## 4 ( 1 )	" "	" "	" "	" "	" "	"*	" "	" "	" "
## 5 ( 1 )	"*	" "	" "	" "	" "	" "	"*	" "	" "
## 6 ( 1 )	"*	" "	" "	" "	" "	" "	"*	" "	" "
## 7 ( 1 )	" "	" "	" "	" "	" "	"*	" "	"*	" "
## 8 ( 1 )	" "	" "	" "	" "	"*	" "	" "	"*	" "
## 9 ( 1 )	"*	" "	" "	" "	" "	" "	"*	" "	" "
## 10 ( 1 )	"*	" "	" "	" "	"*	" "	" "	"*	" "
## 11 ( 1 )	"*	" "	" "	" "	"*	" "	" "	"*	" "
## 12 ( 1 )	"*	" "	" "	" "	"*	" "	" "	"*	" "
## 13 ( 1 )	"*	" "	" "	" "	"*	" "	" "	"*	" "
## 14 ( 1 )	"*	" "	" "	" "	"*	" "	" "	"*	" "
## 15 ( 1 )	"*	" "	" "	" "	"*	" "	" "	"*	" "
## 16 ( 1 )	"*	" "	" "	" "	"*	" "	" "	"*	" "

```

## 17 ( 1 ) "*" " " " " "*" " " "*" " " "*" " "
## 18 ( 1 ) "*" " " " " "*" " " "*" " " "*" " "
## 19 ( 1 ) "*" " " " " "*" " " "*" " " "*" " "
## 20 ( 1 ) "*" "*" " " "*" " " "*" " " "*" " "
## 21 ( 1 ) "*" "*" "*" "*" "*" " " "*" " " "*" " "
## 22 ( 1 ) "*" "*" "*" "*" "*" "*" " " "*" " " "*" " "
## 23 ( 1 ) "*" "*" "*" "*" "*" "*" " " "*" " " "*" " "
## 24 ( 1 ) "*" "*" "*" "*" "*" " " "*" " " "*" " " "*" " "
## 25 ( 1 ) "*" "*" "*" "*" "*" " " "*" " " "*" " " "*" " "
## 26 ( 1 ) "*" "*" "*" "*" "*" "*" " " "*" " " "*" " " "*" " "

## as.factor(rad)3 as.factor(rad)4 as.factor(rad)5 as.factor(rad)6
## 1 ( 1 ) " " " " " " " "
## 2 ( 1 ) " " " " " " " "
## 3 ( 1 ) " " " " " " " "
## 4 ( 1 ) " " " " " " " "
## 5 ( 1 ) " " " " " " " "
## 6 ( 1 ) " " " " " " " "
## 7 ( 1 ) " " " " " " " "
## 8 ( 1 ) " " " " " " " "
## 9 ( 1 ) " " " " " " " "
## 10 ( 1 ) " " " " " " " "
## 11 ( 1 ) " " " " " " " "
## 12 ( 1 ) "*" " " " " " " " "
## 13 ( 1 ) "*" " " " " " " " "
## 14 ( 1 ) "*" " " " " " " " "
## 15 ( 1 ) "*" " " " " " " " "
## 16 ( 1 ) "*" " " " " " " "*" "
## 17 ( 1 ) "*" " " " " "*" " "*" "
## 18 ( 1 ) "*" "*" " " "*" " "*" "
## 19 ( 1 ) "*" "*" " " "*" " "*" "
## 20 ( 1 ) "*" "*" " " "*" " "*" "
## 21 ( 1 ) "*" "*" " " "*" " "*" "
## 22 ( 1 ) "*" "*" " " "*" " "*" "
## 23 ( 1 ) "*" "*" " " "*" " "*" "
## 24 ( 1 ) "*" "*" " " "*" " "*" "
## 25 ( 1 ) "*" "*" " " "*" " "*" "
## 26 ( 1 ) "*" "*" " " "*" " "*" "

## as.factor(rad)7 as.factor(rad)8 as.factor(rad)24 tax ptratio
## 1 ( 1 ) " " " " " " " " " "
## 2 ( 1 ) " " " " " " " " "*" "
## 3 ( 1 ) " " " " " " " " " "
## 4 ( 1 ) " " " " " " " " " "
## 5 ( 1 ) " " " " " " " " " "
## 6 ( 1 ) " " " " " " " " " "
## 7 ( 1 ) " " " " " " " " " "
## 8 ( 1 ) " " " " " " " " " "
## 9 ( 1 ) " " " " " " "*" " " "
## 10 ( 1 ) " " " " " " "*" " " "
## 11 ( 1 ) " " " " " " "*" " " "
## 12 ( 1 ) " " " " " " "*" " " "
## 13 ( 1 ) "*" " " " " "*" " " "*" "
## 14 ( 1 ) "*" " " " " "*" " "*" "
## 15 ( 1 ) "*" "*" " " "*" "*" "
## 16 ( 1 ) "*" "*" " " "*" "*" "

```

```

## 17 ( 1 ) "*" "*" "*" "*"
## 18 ( 1 ) "*" "*" "*" "*"
## 19 ( 1 ) "*" "*" "*" "*"
## 20 ( 1 ) "*" "*" "*" "*"
## 21 ( 1 ) "*" "*" "*" "*"
## 22 ( 1 ) "*" "*" "*" "*"
## 23 ( 1 ) "*" "*" "*" "*"
## 24 ( 1 ) "*" "*" "*" "*"
## 25 ( 1 ) "*" "*" "*" "*"
## 26 ( 1 ) "*" "*" "*" "*"


##          lstat I(tlstat^2) crim:age rm:tax rm:ptratio crim:nox nox:age
## 1 ( 1 ) "*" " " " " " " " " " "
## 2 ( 1 ) "*" " " " " " " " " " "
## 3 ( 1 ) "*" "*" " " " " " " "*" " "
## 4 ( 1 ) "*" " " " " " " "*" "*" " "
## 5 ( 1 ) "*" " " " " " " "*" "*" " "
## 6 ( 1 ) "*" "*" " " " " "*" "*" "*" " "
## 7 ( 1 ) "*" "*" " " " " "*" "*" "*" " "
## 8 ( 1 ) "*" "*" " " " " "*" "*" "*" " "
## 9 ( 1 ) "*" "*" " " " " "*" "*" "*" " "
## 10 ( 1 ) "*" "*" " " " " "*" "*" "*" " "
## 11 ( 1 ) "*" "*" " " " " "*" "*" "*" " "
## 12 ( 1 ) "*" "*" " " " " "*" "*" "*" " "
## 13 ( 1 ) "*" "*" " " " " "*" "*" "*" " "
## 14 ( 1 ) "*" "*" " " " " "*" "*" "*" " "
## 15 ( 1 ) "*" "*" " " " " "*" "*" "*" " "
## 16 ( 1 ) "*" "*" " " " " "*" "*" "*" " "
## 17 ( 1 ) "*" "*" " " " " "*" "*" "*" " "
## 18 ( 1 ) "*" "*" " " " " "*" "*" "*" " "
## 19 ( 1 ) "*" "*" " " " " "*" "*" "*" " "
## 20 ( 1 ) "*" "*" " " " " "*" "*" "*" " "
## 21 ( 1 ) "*" "*" " " " " "*" "*" "*" " "
## 22 ( 1 ) "*" "*" " " " " "*" "*" "*" " "
## 23 ( 1 ) "*" "*" " " " " "*" "*" "*" " "
## 24 ( 1 ) "*" "*" " " " " "*" "*" "*" "*"
## 25 ( 1 ) "*" "*" "*" "*" "*" "*" "*" "*"
## 26 ( 1 ) "*" "*" "*" "*" "*" "*" "*" "*"

##          crim:ptratio
## 1 ( 1 ) " "
## 2 ( 1 ) " "
## 3 ( 1 ) " "
## 4 ( 1 ) " "
## 5 ( 1 ) " "
## 6 ( 1 ) " "
## 7 ( 1 ) "*"
## 8 ( 1 ) "*"
## 9 ( 1 ) " "
## 10 ( 1 ) " "
## 11 ( 1 ) " "
## 12 ( 1 ) " "
## 13 ( 1 ) " "
## 14 ( 1 ) " "
## 15 ( 1 ) " "
## 16 ( 1 ) " "

```

```
## 17 ( 1 ) " "
## 18 ( 1 ) " "
## 19 ( 1 ) " "
## 20 ( 1 ) " "
## 21 ( 1 ) " "
## 22 ( 1 ) " "
## 23 ( 1 ) "* "
## 24 ( 1 ) "* "
## 25 ( 1 ) "* "
## 26 ( 1 ) "* "
```

```
plot(model.subset, scale = "bic")
```



```
neww<- lm(medv~.-chas-rad+as.factor(chas)+as.factor(rad)+I(tlstat)-lstat+I(tlstat^2)-ind
us + rm*tax + rm*ptratio +
      crim*nox + nox*age + crim*ptratio-nox-zn, data = Boston)
summary(neww)
```

```
##
## Call:
## lm(formula = medv ~ . - chas - rad + as.factor(chas) + as.factor(rad) +
##      I(tlstat) - lstat + I(tlstat^2) - indus + rm * tax + rm *
##      ptratio + crim * nox + nox * age + crim * ptratio - nox -
##      zn, data = Boston)
##
## Residuals:
##      Min        1Q    Median        3Q        Max
## -9.988 -2.357 -0.295  1.710 26.452
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   -97.144965   14.593662   -6.657 7.63e-11 ***
## crim           4.877233    2.164966    2.253 0.024720 *
## rm            20.736829    2.146825    9.659 < 2e-16 ***
## age           0.044336    0.027795    1.595 0.111338
## dis          -0.713249    0.145456   -4.904 1.29e-06 ***
## tax           0.039872    0.013576    2.937 0.003473 **
## ptratio       4.358495    0.910488    4.787 2.25e-06 ***
## as.factor(chas)1 3.350686    0.720684    4.649 4.30e-06 ***
## as.factor(rad)2  1.747042    1.220218    1.432 0.152863
## as.factor(rad)3  4.350515    1.106542    3.932 9.67e-05 ***
## as.factor(rad)4  2.251578    0.989912    2.275 0.023372 *
## as.factor(rad)5  2.137704    1.003695    2.130 0.033690 *
## as.factor(rad)6  2.502674    1.221602    2.049 0.041033 *
## as.factor(rad)7  4.025847    1.307510    3.079 0.002195 **
## as.factor(rad)8  2.857172    1.226946    2.329 0.020287 *
## as.factor(rad)24 5.776371    1.495176    3.863 0.000127 ***
## I(tlstat)      -5.487842    0.403798  -13.591 < 2e-16 ***
## I(tlstat^2)     1.268794    0.172725    7.346 8.78e-13 ***
## rm:tax         -0.007640    0.002136   -3.576 0.000384 ***
## rm:ptratio     -0.776008    0.138199   -5.615 3.32e-08 ***
## crim:nox       -1.788827    0.614044   -2.913 0.003743 **
## nox:age        -0.077856    0.046162   -1.687 0.092331 .
## crim:ptratio   -0.190095    0.110191   -1.725 0.085141 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 3.91 on 483 degrees of freedom
## Multiple R-squared:  0.8271, Adjusted R-squared:  0.8193
## F-statistic: 105.1 on 22 and 483 DF,  p-value: < 2.2e-16
```


```
print(paste0("Boston New Model adjusted R squared    ", summary(neww)$adj.r.squared , "
Boston New Model residual standard error    ", summary(neww)$sigma) )
```

```
## [1] "Boston New Model adjusted R squared    0.819268765413949  Boston New Model residu
al standard error    3.90991853940085"
```

```
plot(neww,3, main='Boston New')
```




## Boston New



```
qqnorm(residuals(neww), main='Boston New')
qqline(residuals(neww), col='red')
```

## Boston New

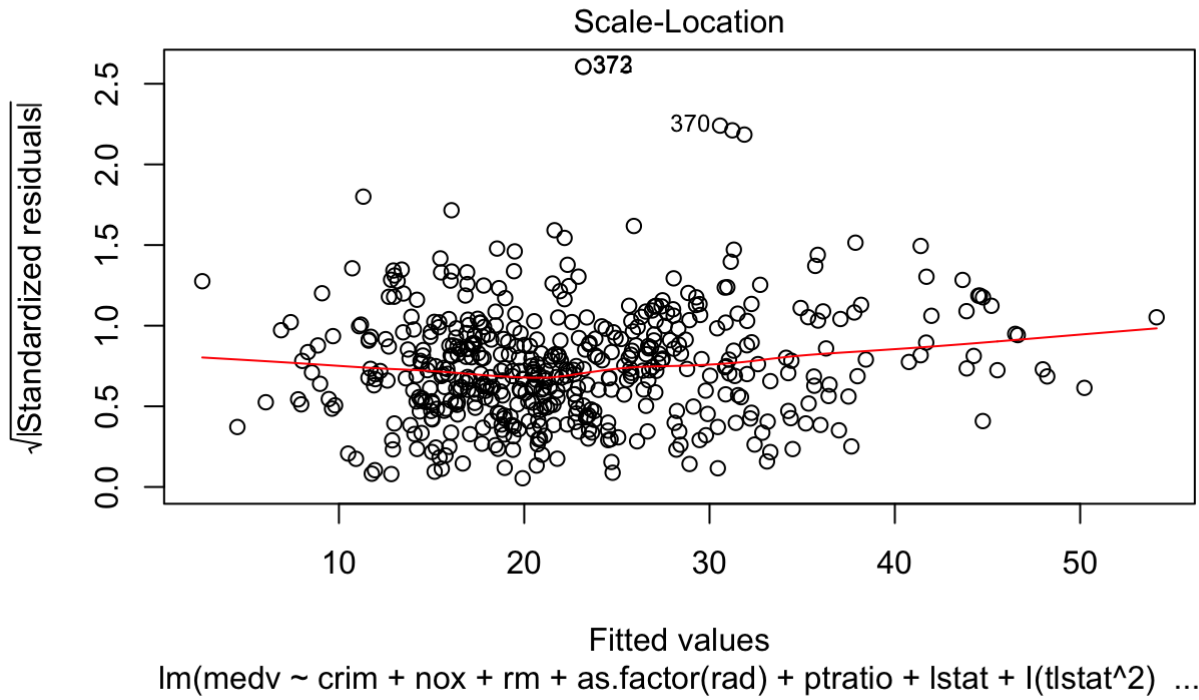


```
print(paste0("Boston Polynomial Model adjusted R squared    ", summary(BestPY)$adj.r.squa
red, "    Boston Polynomial Model residual standard error    ", summary(BestPY)$sigma))
```

```
## [1] "Boston Polynomial Model adjusted R squared    0.811308222633533    Boston Polynomia
l Model residual standard error    3.99509940008713"
```


```
plot(BestPY, 3, main='BestPY')
```

## BestPY



```
qqnorm(residuals(BestPY), main='BestPY')
qqline(residuals(BestPY), col='red')
```

## BestPY



So,

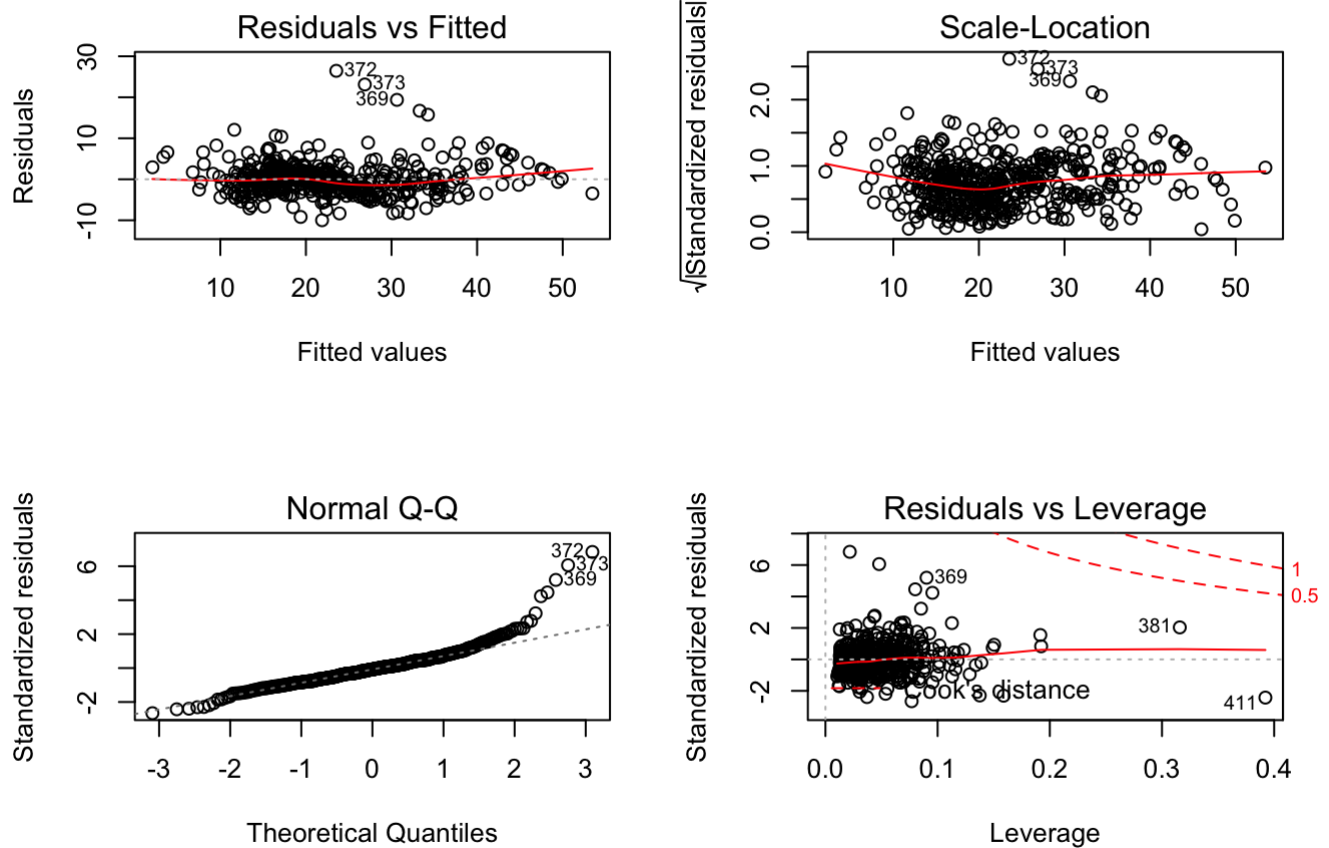
following the KISS principle we will just go with the Boston Polynomial Model as it is the simpler model.

g. **Assesing the model:** Using your model in (f)

- detect 3 points from the data which you think are most probably outliers but not influential points.
- Detect pure leverage points and influential points (if no such points then say not detected, if there are more than 3 then write the most significant 3).
- Calculate the R-Student residuals at the points you find in this part.

Answer goes here:

```
layout(matrix(c(1,2,3,4),2,2))
plot(neww)
```



```
outlierTest(neww)
```

##	rstudent	unadjusted p-value	Bonferroni p
## 372	7.190132	2.4843e-12	1.2571e-09
## 373	6.298073	6.7948e-10	3.4382e-07
## 369	5.339081	1.4417e-07	7.2948e-05
## 370	4.545945	6.9246e-06	3.5038e-03
## 371	4.310024	1.9800e-05	1.0019e-02

I believe that the outliers are points 369, 372, and 373. The most influential points are observations 411 and 381 with a large Cook's distance and the large leverage.

- h. **Collinearity**; Check for multicollinearity in model part (a), part (d), and your model in part (f). Compare the differences in multicollinearity and discuss its possible causes.

Answer goes here:

```
# Evaluate Collinearity
vif(Boston.BM) # variance inflation factors
```

```
##              GVIF Df GVIF^(1/(2*Df))
## crim          1.788890 1          1.337494
## zn             2.495788 1          1.579806
## indus          4.366272 1          2.089563
## as.factor(chas) 1.099046 1          1.048354
## nox            4.653666 1          2.157236
## rm             1.985990 1          1.409252
## age            3.164006 1          1.778765
## dis            4.141903 1          2.035167
## as.factor(rad) 18.629578 8          1.200571
## tax            9.869994 1          3.141655
## ptratio        2.241516 1          1.497169
## lstat          2.922144 1          1.709428
```

```
sqrt(vif(Boston.BM)) > 2 # problem?
```

```
##              GVIF      Df GVIF^(1/(2*Df))
## crim          FALSE FALSE             FALSE
## zn             FALSE FALSE             FALSE
## indus          TRUE  FALSE             FALSE
## as.factor(chas) FALSE FALSE             FALSE
## nox            TRUE  FALSE             FALSE
## rm             FALSE FALSE             FALSE
## age            FALSE FALSE             FALSE
## dis            TRUE  FALSE             FALSE
## as.factor(rad)  TRUE   TRUE             FALSE
## tax            TRUE  FALSE             FALSE
## ptratio        FALSE FALSE             FALSE
## lstat          FALSE FALSE             FALSE
```

```
vif(BestPY) # variance inflation factors
```

```
##              GVIF Df GVIF^(1/(2*Df))
## crim          12239.606245 1          110.632754
## nox            47.130197 1           6.865144
## rm             75.579221 1           8.693631
## as.factor(rad)  30.654748 8           1.238529
## ptratio        129.795219 1          11.392770
## lstat          5.490035 1           2.343082
## I(tlstat^2)     2.564882 1           1.601525
## I(tcrim^2)       8.469944 1           2.910317
## age            82.609374 1           9.088970
## tax           172.425690 1          13.131096
## indus           4.358181 1           2.087626
## crim:age        213.762479 1          14.620618
## rm:tax          152.238138 1          12.338482
## rm:ptratio      145.439929 1          12.059848
## crim:nox        426.080019 1          20.641706
## nox:age         218.259411 1          14.773605
## crim:ptratio    12986.607978 1        113.958799
```

```
sqrt(vif(BestPY)) > 2 # problem?
```

```
##              GVIF      Df GVIF^(1/(2*Df))
## crim          TRUE FALSE              TRUE
## nox            TRUE FALSE              TRUE
## rm             TRUE FALSE              TRUE
## as.factor(rad) TRUE  TRUE              FALSE
## ptratio        TRUE FALSE              TRUE
## lstat          TRUE FALSE              FALSE
## I(tlstat^2)    FALSE FALSE              FALSE
## I(tcrim^2)     TRUE FALSE              FALSE
## age            TRUE FALSE              TRUE
## tax            TRUE FALSE              TRUE
## indus          TRUE FALSE              FALSE
## crim:age       TRUE FALSE              TRUE
## rm:tax         TRUE FALSE              TRUE
## rm:ptratio     TRUE FALSE              TRUE
## crim:nox       TRUE FALSE              TRUE
## nox:age        TRUE FALSE              TRUE
## crim:ptratio   TRUE FALSE              TRUE
```

```
vif(neww) # variance inflation factors
```

```
##              GVIF Df GVIF^(1/(2*Df))
## crim          11455.416198 1      107.029978
## rm             75.160001 1       8.669487
## age            20.221256 1       4.496805
## dis            3.098956 1       1.760385
## tax            172.940773 1      13.150695
## ptratio        128.350853 1      11.329204
## as.factor(chas) 1.106859 1       1.052074
## as.factor(rad)  19.817951 8       1.205220
## I(tlstat)       5.386236 1       2.320826
## I(tlstat^2)     2.435865 1       1.560726
## rm:tax          152.120733 1      12.333723
## rm:ptratio      144.537280 1      12.022366
## crim:nox        420.783738 1      20.513014
## nox:age         33.349899 1       5.774937
## crim:ptratio    12138.310514 1     110.174001
```

```
sqrt(vif(neww)) > 2 # problem?
```

```
##          GVIF      Df GVIF^(1/(2*Df))
## crim          TRUE FALSE          TRUE
## rm            TRUE FALSE          TRUE
## age           TRUE FALSE          TRUE
## dis           FALSE FALSE          FALSE
## tax           TRUE FALSE          TRUE
## ptratio       TRUE FALSE          TRUE
## as.factor(chas) FALSE FALSE          FALSE
## as.factor(rad)  TRUE  TRUE          FALSE
## I(tlstat)       TRUE FALSE          FALSE
## I(tlstat^2)     FALSE FALSE          FALSE
## rm:tax          TRUE FALSE          TRUE
## rm:ptratio      TRUE FALSE          TRUE
## crim:nox        TRUE FALSE          TRUE
## nox:age         TRUE FALSE          TRUE
## crim:ptratio    TRUE FALSE          TRUE
```

As we can see from above there are some issues with multicollinearity. These issues may be due to having an  $R_k^2$  of  $\geq 0.9$  or even just from interaction terms.